


Phase 1 — Multi-Query Decomposition

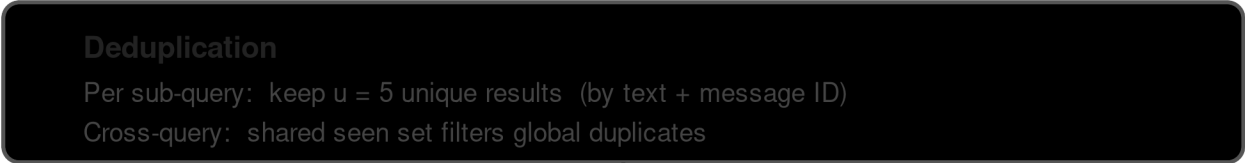
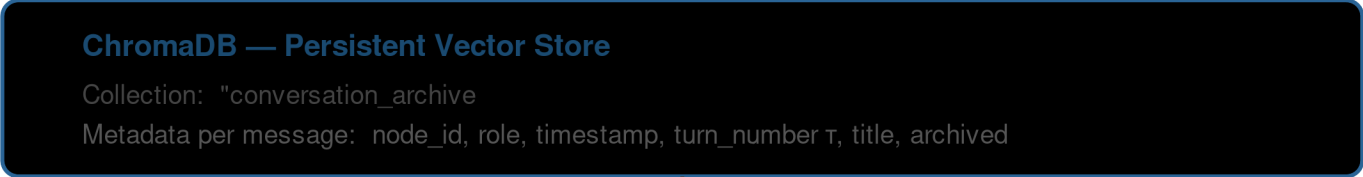


Each sub-query dispatched in parallel to the vector store

Phase 2 — Embedding and Vector Search

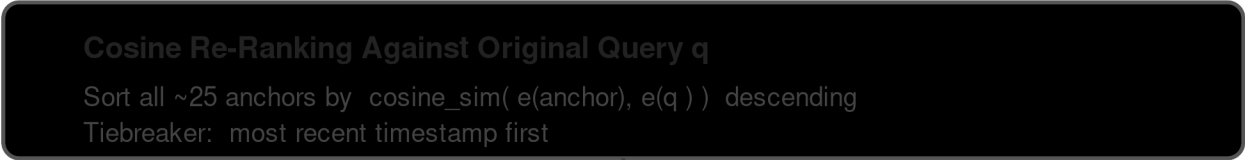


Index Path
add_message()

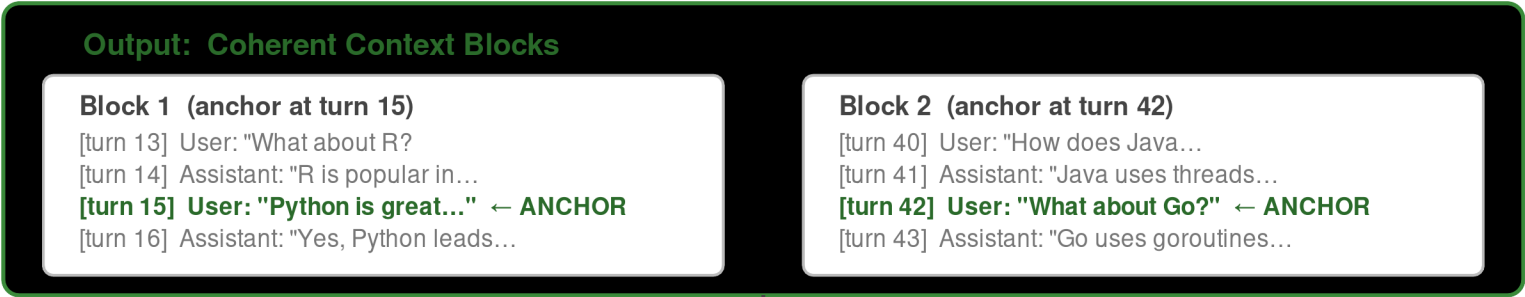
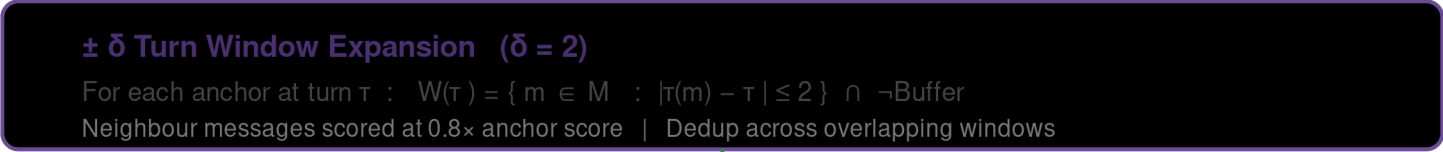


Result: $R = R$ (up to $n \times u \approx 25$ unique anchor messages)

Phase 3 — Re-Rank Anchors



Phase 4 — Turn-Based Context Expansion



$R = k$ coherent blocks → injected into LLM prompt